



Student Depression Classification Pipeline

ML Classification System with MLOps, API Serving, Monitoring,
CI/CD & Deployment

Course: MTA — Advanced Analytics II

Author: Enosh

Date: February 12, 2026

GitHub: [enosh729-design/Student-Depression-Predictor](https://github.com/enosh729-design/Student-Depression-Predictor)

Tech Stack: Python · scikit-learn · FastAPI · Streamlit · Docker · Prometheus · Grafana ·
W&B · Render

Table of Contents

1. Introduction & Problem Statement
2. Architecture & Pipeline Diagram
3. Data Layer
4. Model Training & Experimentation
5. Model Performance & Best Hyperparameters
6. W&B Experiment Tracking
7. Backend API (FastAPI)
8. Containerization & Monitoring
9. Frontend (Streamlit)
10. Testing & Code Quality
11. CI/CD & Deployment
12. Business Value
13. Conclusion

1. Introduction & Problem Statement


Student mental health is a pressing concern across higher education institutions globally. Depression among university students leads to lower academic performance, increased dropout rates, higher healthcare costs, and long-term career impacts. Early identification of at-risk students through data-driven approaches can enable proactive intervention and substantially improve student outcomes.


This project builds a **complete binary classification system** to predict student depression risk based on easily collectible lifestyle factors. The solution follows modern **MLOps best practices** and demonstrates the entire machine learning lifecycle — from data ingestion and versioning to model training, API serving, monitoring, CI/CD, and cloud deployment.

Objectives

- Store and load data from a serverless PostgreSQL database (Neon Postgres)
- Build a scikit-learn pipeline with preprocessing and classification
- Perform hyperparameter tuning and track experiments with Weights & Biases
- Serve predictions via a FastAPI REST API with Prometheus metrics
- Build an interactive Streamlit frontend
- Containerize with Docker, monitor with Prometheus + Grafana
- Implement CI/CD with GitHub Actions, deploy to Render

2. Architecture & Pipeline Diagram





3. Data Layer

3.1 Dataset Overview

The dataset contains **100,000 student records** with lifestyle factors and a binary depression label. The data was pre-collected and provided as a CSV file.

Feature	Type	Description
Student_ID	int	Unique identifier
Age	int	Student age (18–24)
Gender	str	Male / Female
Department	str	Science, Engineering, Medical, Arts, Business
CGPA	float	Cumulative GPA (0.0–4.0)
Sleep_Duration	float	Daily sleep hours
Study_Hours	float	Daily study hours
Social_Media_Hours	float	Daily social media usage
Physical_Activity	int	Physical activity score (0–150)
Stress_Level	int	Stress level (0–10)
Depression	bool	Target: True (10.06%) / False (89.94%)

⚠ Class Imbalance: Only 10.06% of records are labeled as "Depression". This imbalance was addressed using `class_weight="balanced"` in the Random Forest classifier.

3.2 Neon Postgres Integration

The raw CSV was uploaded to a **Neon Postgres** (serverless PostgreSQL) database using the `data/load_to_postgres.py` script. The training pipeline loads data from Postgres via SQLAlchemy with a CSV fallback:

```
# From Postgres (primary)
from src.data_loader import load_data_from_postgres
df = load_data_from_postgres()

# From CSV (fallback)
from src.data_loader import load_data_from_csv
df = load_data_from_csv("data/student_lifestyle_100k.csv")
```

4. Model Training & Experimentation

4.1 Preprocessing Pipeline

A `ColumnTransformer` handles two types of features:

- **Numerical (7 features):** Age, CGPA, Sleep_Duration, Study_Hours, Social_Media_Hours, Physical_Activity, Stress_Level → `StandardScaler`
- **Categorical (2 features):** Gender, Department → `OneHotEncoder` (with `handle_unknown="ignore"`)

4.2 Classification Model

A `RandomForestClassifier` was chosen for its robustness, interpretability, and strong performance on tabular data. The preprocessor and classifier are wrapped in a single `sklearn.Pipeline` for clean serialization and deployment.

4.3 Hyperparameter Tuning

`RandomizedSearchCV` was used with the following search space:

Parameter	Search Space
<code>n_estimators</code>	[50, 100, 200, 300, 500]
<code>max_depth</code>	[5, 10, 15, 20, None]
<code>min_samples_split</code>	[2, 5, 10, 20]
<code>min_samples_leaf</code>	[1, 2, 4, 8]
<code>max_features</code>	["sqrt", "log2", None]
<code>class_weight</code>	["balanced", "balanced_subsample"]

- **Iterations:** 50 random combinations
- **Cross-validation:** 3-fold stratified
- **Scoring metric:** F1-score (optimizes for precision–recall balance)
- **Total fits:** 150 (50 × 3 folds)

5. Model Performance & Best Hyperparameters

5.1 Best Hyperparameters

Parameter	Best Value
n_estimators	100
max_depth	5
min_samples_split	20
min_samples_leaf	4
max_features	sqrt
class_weight	balanced

5.2 Test Set Metrics

Metric	Score
Accuracy	0.7356
F1-Score	0.3226
ROC-AUC	0.7027
Precision	0.2173
Recall	0.6257
Best CV F1	0.3273

Key Insight: The `class_weight="balanced"` setting was critical. Without it, the model achieved 89.9% accuracy but only 1.29% recall — it predicted "No Depression" for nearly every input. With balanced weights, recall improved to 62.57%, enabling meaningful depression detection at the cost of some accuracy.

6. W&B Experiment Tracking

All experiments are tracked in **Weights & Biases** under the project `student-depression-classifier`. The training pipeline logs:

- **Metrics:** Accuracy, F1, ROC-AUC, Precision, Recall
- **Visualizations:** Confusion Matrix, ROC Curve
- **Parameters:** All hyperparameters from RandomizedSearchCV
- **Artifacts:** Best model pipeline (`.joblib`) + metrics JSON




Figure 1: W&B Experiment Dashboard




Figure 2: Confusion Matrix




Figure 3: ROC Curve

Model Artifact

The best model is saved as `models/best_model.joblib` (67 MB) and also registered as a W&B artifact named `best-model` with type `model`. This enables versioning, lineage tracking, and reproducibility.

7. Backend API (FastAPI)

7.1 Endpoints

Method	Path	Description
POST	/predict	Submit student data, returns depression prediction with probabilities
GET	/health	Health check — returns API status and model loaded status
GET	/metrics	Prometheus-format metrics for monitoring
GET	/docs	Auto-generated Swagger UI documentation

7.2 Pydantic Validation

All request/response payloads are validated with Pydantic models. For example, `StudentInput` enforces field-level constraints: `Age` must be 15–30, `CGPA` must be 0.0–4.0, `Stress_Level` must be 0–10, etc. Invalid inputs return HTTP 422 with detailed error messages.

7.3 Postman Testing




Figure 4: Successful Prediction (200 OK)

The screenshot shows the Postman interface with a POST request to `https://student-depression-predictor-c3ce.onrender.com/predict`. The request body is currently empty. The response status is `422 Unprocessable Entity`, with a timestamp of `1.98 s` and a size of `534 B`. The response body is a JSON object:

```
1: {  
2:   "detail": [  
3:     {  
4:       "type": "missing",  
5:       "loc": [  
6:         "body",  
7:         "Age"  
8:       ],  
9:       "msg": "Field required",  
10:      "input": {}  
11:    }  
12:  ]  
13:}
```

Figure 5: Validation Error (422 Unprocessable Entity)

The screenshot shows the Postman interface with a GET request to `https://student-depression-predictor-c3ce.onrender.com/health`. The request body is empty. The response status is `200 OK`, with a timestamp of `1.01 s` and a size of `401 B`. The response body is a JSON object:

```
1: {  
2:   "status": "healthy",  
3:   "model loaded": true,  
4:   "version": "1.0.0"  
5: }
```

Figure 6: Health Check (200 OK)

8. Containerization & Monitoring

8.1 Docker Setup

The FastAPI backend runs in a Docker container built from `python:3.11-slim`. A `docker-compose.yml` orchestrates three services:

Service	Image	Port
FastAPI API	custom (Dockerfile)	8000
Prometheus	prom/prometheus:latest	9090
Grafana	grafana/grafana:latest	3000

8.2 Prometheus Metrics

The API exposes custom Prometheus metrics via `/metrics`:

- `prediction_requests_total` — Counter with status labels (success/error)
- `prediction_latency_seconds` — Histogram with configurable buckets
- `prediction_results_total` — Counter with outcome labels (Depression/No Depression)
- `model_loaded` — Gauge (1 = loaded, 0 = not loaded)

8.3 Grafana Dashboards (6 Panels)

A pre-provisioned Grafana dashboard (`api_dashboard.json`) provides real-time visibility:

1. **Request Count Rate** — Prediction requests per second over time
2. **Latency Percentiles** — p50, p95, p99 response latency
3. **Prediction Outcomes** — Depression vs No Depression rates
4. **Model Status** — UP/DOWN indicator
5. **Total Predictions** — Cumulative prediction count
6. **Error Rate** — Percentage of failed requests




Figure 7: Grafana Monitoring Dashboard

9. Frontend (Streamlit)

An interactive Streamlit UI at `http://localhost:8501` provides a user-friendly interface for making predictions:

- **Input Form:** Sliders and dropdowns for all 9 features (Age, Gender, Department, CGPA, Sleep, Study Hours, Social Media, Physical Activity, Stress Level)
- **API Health Check:** Sidebar indicator showing API connection status
- **Prediction Display:** Color-coded result box with confidence scores
- **Detailed Probabilities:** Breakdown of No Depression / Depression probabilities

The screenshot shows the Streamlit interface for the Student Depression Risk Predictor. On the left, there's a sidebar with 'System Status' (API Online), 'Model: Loaded Version: 1.0.0', and an 'About' section. The main area has a dark blue header with 'AI-POWERED ASSESSMENT', 'Student Depression Risk Predictor', and a subtext 'Predict depression risk by analyzing lifestyle data. Machine learning model trained on 100,000 student records.' Below the header is a section titled 'Enter Student Details' with three tabs: 'DEMOGRAPHICS', 'ACADEMICS', and 'LIFESTYLE'. Under 'DEMOGRAPHICS', there are sliders for Age (set to 20), CGPA (set to 3.00), and Social Media (hrs/day) (set to 3.00). Under 'ACADEMICS', there are dropdowns for Gender (Male) and Department (Science), and sliders for Study Hours / day (set to 4.00) and Sleep Duration (hrs) (set to 7.00). Under 'LIFESTYLE', there are sliders for Physical Activity (set to 80) and Stress Level (set to 5). At the bottom is a large blue button labeled 'Analyze Depression Risk'.

Figure 8: Streamlit Interface — Input Form




Figure 9: Streamlit Interface — Prediction Result

10. Testing & Code Quality

10.1 Unit Tests (pytest)

15 tests across 3 test files — all passing:

Test File	Tests	Covers
test_data_validation.py	5	Schema validation, missing columns, feature/target split, target dtype, empty DataFrame
test_model_inference.py	5	Prediction dict format, valid prediction values, probability sum, batch prediction, label–prediction match
test_api_endpoints.py	5	GET /health, POST /predict (valid), POST /predict (missing fields → 422), POST /predict (invalid range → 422), GET /metrics

```
$ python -m pytest tests/ -v --tb=short
=====
 15 passed, 1 warning in 1.36s =====
```

10.2 Linting Scores

Tool	Result
Flake8	0 errors, 0 warnings ✓
Pylint	Minor non-critical notes only (R0914 too-many-locals in train function) ✓

11. CI/CD & Deployment

11.1 GitHub Actions

Two CI/CD workflows are configured in `.github/workflows/`, both triggered on push to `main`:

Backend Pipeline (`backend.yml`)

```
push to main —> Lint (Flake8 + Pylint) —> Test (pytest) —> Deploy to Render
```

Frontend Pipeline (`frontend.yml`)

```
push to main —> Lint (Flake8 + Pylint) —> Test (pytest) —> Deploy to Render
```

11.2 Render Deployment

Both services are deployed to **Render** cloud platform:

Service	Type	Live URL
FastAPI Backend	Docker	student-depression-predictor-c3ce.onrender.com
Streamlit Frontend	Python	student-depression-frontend.onrender.com

11.3 GitHub Repository

All code is hosted on a public GitHub repository:

<https://github.com/enosh729-design/Student-Depression-Predictor>

12. Business Value

12.1 Problem

Student mental health is a critical and growing concern. Depression leads to lower academic performance, higher dropout rates, increased healthcare costs, and long-term career impacts. Traditional screening methods are manual, expensive, and don't scale.

12.2 Solution

This ML system enables **early identification** of at-risk students using easily collectible lifestyle data. Key value propositions:

1. **Proactive Intervention** — Flag at-risk students before grades drop, enabling counselors to act early
2. **Resource Optimization** — Direct limited counseling resources to students with the highest need
3. **Data-Driven Policy** — Provide administrators with evidence-based insights for wellness program design
4. **Scalable Screening** — Process thousands of students instantly vs. weeks of manual assessment
5. **Privacy-Preserving** — Uses lifestyle metrics (sleep, stress, activity) rather than sensitive medical records

12.3 ROI Estimate

A university with 10,000 students and a 15% depression rate can save an estimated **\$2–5M annually** through reduced dropout costs and optimized counseling. Early intervention has been shown to improve retention rates by **10–15%**.

13. Conclusion

This project demonstrates a **complete, production-grade ML classification system** following modern MLOps best practices. The solution covers the full lifecycle:

- **Data Layer:** Neon Postgres with SQLAlchemy data loading
- **Training:** scikit-learn pipeline with RandomizedSearchCV and W&B tracking
- **Serving:** FastAPI REST API with Pydantic validation and Prometheus metrics
- **Monitoring:** Docker Compose with Prometheus + Grafana (6 dashboards)
- **Frontend:** Interactive Streamlit UI
- **Quality:** 15 pytest tests, Flake8/Pylint clean code
- **CI/CD:** GitHub Actions (lint → test → deploy) for both backend and frontend
- **Deployment:** Render cloud (FastAPI + Streamlit)

The key technical insight was the importance of handling class imbalance — using `class_weight="balanced"` and F1 scoring improved recall from 1.29% to 62.57%, transforming the model from a trivial majority-class predictor into a useful depression screening tool.

The system is designed for real-world adoption by university counseling departments, with a simple web interface for non-technical users and robust API endpoints for integration with existing student information systems.

Demo Video

 **Demo Video Link:** [INSERT 5-MINUTE VIDEO LINK HERE]

The video walks through: local setup → data loading → model training → API testing → Streamlit UI → Grafana dashboards → Render deployment.

Project Links

Resource	URL
GitHub Repository	github.com/enosh729-design/Student-Depression-Predictor
Live API	student-depression-predictor-c3ce.onrender.com/docs
Live Frontend	student-depression-frontend.onrender.com
W&B Project	wandb.ai/YOUR_PROJECT

